

# Product Requirements Document: Project GQE-MTS

TEAM QAT

February 1, 2026

## Hybrid Generative Quantum-Enhanced Memetic Tabu Search for LABS

**Team Name:** QAT

**GitHub Repository:** <https://github.com/rnsln/iqhahack2026-nvidia>

**Github Profiles:**

Hatice Boyar: <https://github.com/dhaticeboyar>

Eren Aslan: <https://github.com/rnsln>

Hüseyin Umut Işık: <https://github.com/HUmutI>

**Web Page:** COMING SOON...

## 1 Team Roles & Responsibilities (PICs)

- **Project Lead (Architect): HATICE BOYAR and CHANG JEN YU**  
Responsible for the algorithmic "North Star," high-level system integration, and defining the Hamiltonian. Ensures the theoretical validity of the ansatz.
- **GPU Acceleration PIC (Builder): HUSEYIN UMUT ISIK**  
Lead for the CUDA-Q 0.13.0 implementation and the CuPy-based parallel evaluation engine. Responsible for optimizing memory coalescing in CUDA kernels and "Zombie Instance" prevention on Brev.dev.
- **Quality Assurance PIC (Verifier): ILAYDA DILEK**  
Owner of the Verification Strategy. Responsible for the automated unit test suite and Python library **Hypothesis** to guard against AI hallucinations and ensure physical consistency.
- **Technical Marketing PIC (Storyteller): EREN ASLAN**  
The data analyst. Responsible for converting raw logs into performance visualizations (Speedup/Convergence plots). **Key Initiative:** Developing a **Live Interactive Web Dashboard** (Node.js + Express.js + HTML). This platform will feature

dynamic charts showing the real-time descent of the MTS algorithm and 3D visualizations of the Quantum Energy Landscape, providing an immersive storytelling experience for the judges.

## 2 The Architecture: Hybrid GQE-MTS

**Owner:** Project Lead - CHANG JEN YU and Hatice Boyar

### Methodology: Hybrid GQE-MTS Architecture

Our approach synergizes the **global exploration capabilities of Quantum Computing (GQE)** with the **local exploitation capabilities of Classical Heuristic Algorithms (MTS)**[2] to solve the highly complex Low Autocorrelation Binary Sequences (LABS)[3] problem.

### Evolution and Rationale

Initially, our research utilized a GPU-accelerated QE-MTS approach. However, we observed that this method occasionally yielded inconsistent results due to stochastic errors. While we considered variational methods such as VQE[4] or QAOA[5] as alternatives, these approaches suffer significantly from the "**Barren Plateau**" problem[6] in high-dimensional optimization landscapes.

Consequently, we pivoted to **Generative Quantum Eigensolvers (GQE)**[1] augmented by **Generative AI** for circuit synthesis. This approach avoids the pitfalls of traditional variational parameter optimization by guiding the circuit generation process.

### Key Innovations

We introduce two distinct novelties to adapt GQE for the LABS problem:

#### 1. Transfer Learning[7] for GQE Acceleration:

Standard GQE requires substantial computational resources. We integrated **Transfer Learning** to accelerate prediction. By training on smaller sub-problems and transferring the geometric insights to larger instances, we significantly reduce the training latency.

#### 2. Non-Chemical Domain Adaptation:

Unlike standard GQE rooted in quantum chemistry, we treat LABS as a distribution problem. We construct our sampling space using **2-body and 4-body operators** derived from quantum annealing principles.

The methodology is divided into three distinct phases:

#### Phase 1: Quantum Generative Optimization (GQE Phase)

The goal of this phase is not to find the perfect solution immediately, but to **locate the "Basin of Attraction"** where the good solutions reside.

1. **Physical Modeling:** We map the LABS autocorrelation energy into a Quantum Hamiltonian ( $H$ ).

2. **Operator Pool Design:**

- We utilize problem-specific operators (two-body and four-body terms) derived from the interaction graph.
- **Hierarchical Parameters:** We employ a decaying parameter strategy ( $\pi \rightarrow \pi/8$ ). This allows the GQE to perform large geometric rotations (coarse tuning) first, followed by fine quantum interference adjustments (fine tuning).

3. **Transfer Learning Strategy (Input/Output Adaptation):**

To solve the target problem size ( $N_{target}$ ), we employ a specific training protocol on smaller instances ( $N_{train}$ ):

- **Zero-Padding Input:** We define the model architecture based on the maximum target size ( $N_{target}$ ). When training on smaller sub-problems ( $N_{train} < N_{target}$ ), the input vectors are **zero-padded** to match the dimensionality of  $N_{target}$ .
- **Output Truncation:** During the training process, the model generates a circuit for the full dimension. However, for loss calculation, we **truncate or mask the output** (the extra qubits/gates beyond  $N_{train}$ ) to evaluate the energy only within the valid subspace of the current training instance.
- This allows the generative model to learn the fundamental geometric features of LABS interactions on small  $N$  and transfer these weights to larger  $N$  without retraining from scratch.

4. **GQE Engine Execution:**

- The engine samples operators from the pool using the pre-trained/transferred model.
- It minimizes the energy expectation value  $\langle H \rangle$ .
- **Result:** An optimized **Quantum Ansatz** where the wavefunction is concentrated on low-energy states.

## Phase 2: Bridging & Filtering Phase

This phase acts as the interface between the quantum and classical worlds, converting "quantum probabilities" into "high-quality seeds."

1. **Sampling:** We execute 1000 shots on the optimized circuit to obtain binary sequences (Bitstrings).
2. **Energy Validation:** Convert bitstrings back to the spin format and use a classical CPU to calculate the exact LABS energy.

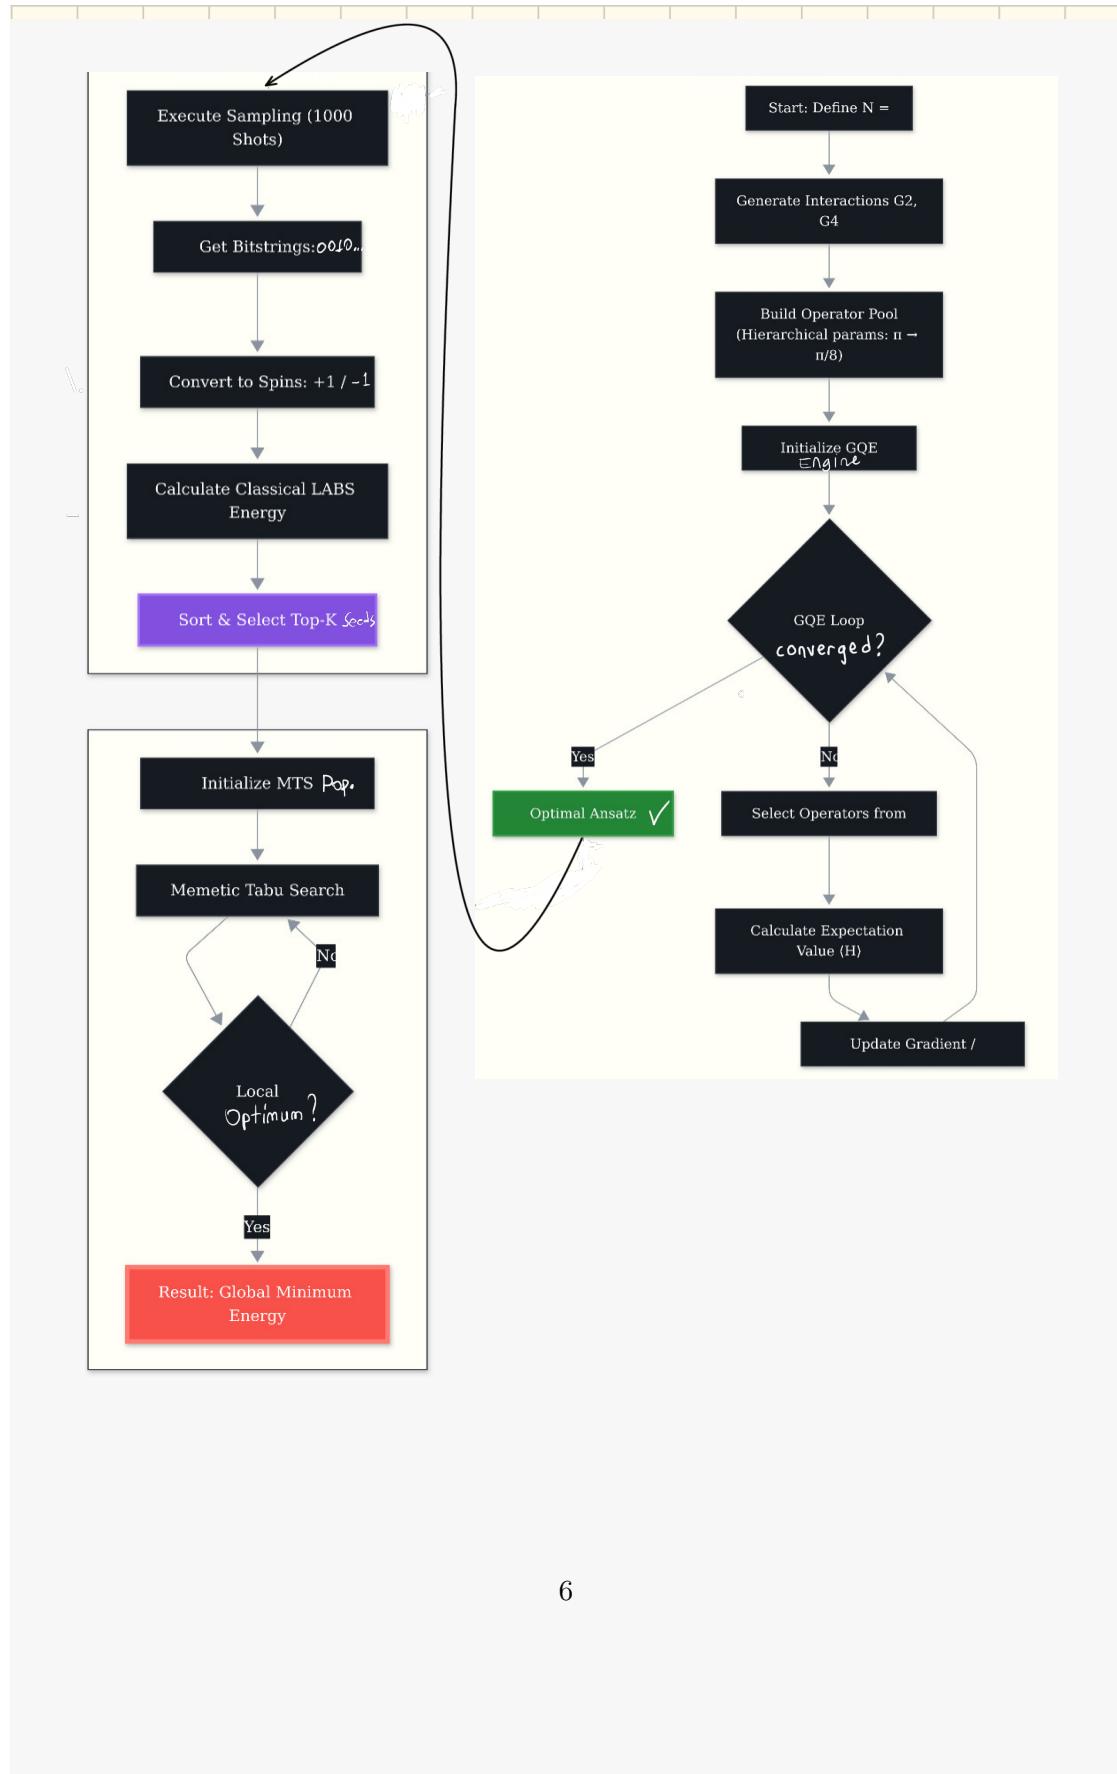
3. **Elitism Filtering:** From the 1000 samples, we select only the top  $K$  (e.g., 20) unique sequences with the lowest energy. These become our "**Golden Seeds.**"

### Phase 3: Classical Memetic Search (MTS Phase)

This phase leverages cost-effective classical computing to perform the final optimization sprint.

1. **Seed Injection:** The "Golden Seeds" are injected as the **Initial Population** for the MTS.
2. **Tabu Search:** The MTS performs bit-flips within the neighborhood of these seeds, utilizing a Tabu list to prevent cycling back to previous states.
3. **Convergence:** Since the starting points are already deep within the energy valleys, the MTS can slide down to the **Global Minimum** extremely fast.

## System Architecture Flowchart



## Why This Method Works

### 1. Winning at the Starting Line:

- Pure random search is like dropping onto the Himalayas blindly; it is statistically unlikely to land near the peak of Mt. Everest.
- Our method uses Quantum Radar (GQE) to scan for "where the mountain is" and drops the climbing team (MTS) halfway up the slope.

### 2. Physics-Aware:

- The Operator Pool is not random; it is strictly defined by the structure of the LABS problem. This means every step the quantum circuit takes is a valid physical evolution, making it far more efficient than blind parameter optimization.

### 3. Fault Tolerance:

- Even if the GQE does not find the perfect Ground State, as long as it finds a sufficiently low excited state, the MTS can easily complete the "last mile" of optimization.

## 3 The Acceleration Strategy

**Owner:** GPU Acceleration PIC - HUSEYIN UMUT ISIK

### Quantum Acceleration (CUDA-Q)

- **Strategy:** We utilize the `nvidia` backend in CUDA-Q 0.13.0 for high-performance state-vector simulation.
  - **Multi-GPU Scaling:** For sequence lengths  $N > 30$ , we leverage `nvidia-mgpu` on Brev to distribute the large state-vector ( $2^N$  amplitudes) across multiple L4/A100 GPUs using MPI decomposition.
  - **Circuit Optimization:** We use `cudaq.optimize` passes to fuse single-qubit gates and reduce circuit depth before execution.

### Classical Acceleration (MTS)

- **Strategy:** The bottleneck of MTS is evaluating the  $N$  single-flip neighbors.
  - **Delta-Evaluation Optimization:** Instead of recalculating energy from scratch ( $O(N^2)$ ), we compute the change in energy ( $\Delta E$ ) for a bit-flip, which reduces complexity to  $O(N)$ .

- **CuPy Parallelization:** We implement a custom CUDA kernel via CuPy to evaluate all possible  $N$  flips and Tabu criteria in parallel, achieving a massive speedup over serial CPU iterators. We tune GPU performance by adjusting kernel granularity (threads per block and grid dimensions) to maximize occupancy and minimize launch overhead.

- **Hardware Targets:**

- **Dev:** qBraid (CPU) for logic and unit testing.
- **Production:** Brev (A100-80GB) for final benchmarks on bigger  $N$  values with different hyperparameters like FP32, FP64, MGPU on/off, etc.

## 4 The Verification Plan

**Owner:** Quality Assurance PIC - ILAYDA DILEK

### Core Correctness Checks

- **Check 1 (Symmetry & Invariants):**
  - **Reversal Symmetry:**  $E(S) = E(S_{\text{reversed}})$ .
  - **Negation Symmetry:**  $E(S) = E(-S)$ .
  - **Energy Bounds:** Ensure computed energy does not violate theoretical lower bounds for LABS.
- **Check 2 (Ground Truth Calibration):**
  - **Small N:** For  $N = 7$ , must return  $E = 3$ .
  - **Medium N:** For  $N = 20$ , must target known minimum  $E = 16$  (Merit Factor  $\approx 6.0$ ).
  - Verify results against known Barker Codes for applicable lengths.
- **Check 3 (Operator Pool & Gradients):**
  - Unit tests to ensure the 4-qubit  $R_{ZZ}$  chains in the DCQO circuit correctly preserve parity.
  - Finite-difference checks to verify analytic gradients used in the GQE loop.

### AI Hallucination Guardrails

- We use a "Test-Driven Prompting" strategy: AI agents are given the test file before the implementation request. Any code that does not pass the automated `tests.py` suite is immediately rejected for refactoring.

## 5 Execution Strategy & Success Metrics

**Owner:** Technical Marketing PIC - EREN ASLAN

### Agentic Workflow

- **Orchestration:**
  - **Agent Alpha (Architect):** Decomposes tasks into atomic issues (e.g., "Implement Interaction Graph Generation").
  - **Agent Gamma (QA):** Reviews PRs, analyzes `pytest` logs, and spots "silent failures" where code runs but produces physics-violating results.
- **Vibe Coding Tip:** We will keep a live "Vibe Check" log to record where the AI agents struggled with CUDA-Q syntax vs. classical Python logic.

### Success Metrics

- **Target Hit Rate:** Achieve the global minimum for  $N = 27$  in under 120 seconds.
- **Speedup:** Target a 20x speedup on the Brev A100 vs. the qBraid CPU baseline.
- **Approximation Ratio:** Maintain an energy ratio  $> 0.9$  relative to best-known values for  $N = 40$ .

## 6 Resource Management Plan

**Owner:** GPU Acceleration PIC

- **Credit Allocation Strategy (Brev.dev):**
  - **Prototyping Phase:** Initial development, debugging, and unit testing are performed on cost-efficient GPU instances using small-to-moderate problem sizes. During this phase, all code changes are validated through the automated `tests.py` suite to prevent unnecessary GPU usage caused by incorrect or unstable implementations.
  - **Benchmarking Phase:** High-performance GPU instances are reserved exclusively for large-scale production runs and benchmarking experiments (e.g.,  $N \geq 30$ ), after functional correctness and stability have been verified.
  - **Operational Buffer:** A portion of available credits is intentionally reserved to accommodate re-runs, profiling, debugging, and unexpected operational overhead.

- **Zombie Prevention and Resource Guardrails:**

- **Automated Monitoring:** A lightweight background script monitors GPU utilization via `nvidia-smi` at regular intervals to detect idle or stalled instances.
- **Policy Enforcement:** If sustained low utilization is detected over a pre-defined time window, the instance is automatically terminated to prevent unintended credit consumption.

## References

- [1] Nakaji, Kouhei, et al. *The generative quantum eigensolver (GQE) and its application for ground state search*. arXiv preprint arXiv:2401.09253 (2024).
- [2] Gomez Cadavid, A., Chandarana, P., Romero, S. V., et al. *Scaling advantage with quantum-enhanced memetic tabu search for LABS*. arXiv preprint arXiv:2511.04553 (2025).
- [3] Packebusch, T., and Mertens, S. *Low autocorrelation binary sequences*. Journal of Physics A: Mathematical and Theoretical, 49(16), 165001 (2016). DOI: 10.1088/1751-8113/49/16/165001.
- [4] Tilly, J., Chen, H., Cao, S., et al. *The Variational Quantum Eigensolver: A review of methods and best practices*. Physics Reports, 986, 1–128 (2022). DOI: 10.1016/j.physrep.2022.08.003.
- [5] Farhi, E., Goldstone, J., and Gutmann, S. *A Quantum Approximate Optimization Algorithm*. arXiv preprint arXiv:1411.4028 (2014). URL: <https://arxiv.org/abs/1411.4028>.
- [6] Larocca, M., Thanasilp, S., Wang, S., et al. *Barren plateaus in variational quantum computing*. Nature Reviews Physics, 7(4), 174–189 (2025). DOI: 10.1038/s42254-025-00813-9.
- [7] Zhuang, F., Qi, Z., Duan, K., et al. *A Comprehensive Survey on Transfer Learning*. Proceedings of the IEEE, 109(1), 43–76 (2020). arXiv preprint arXiv:1911.02685.